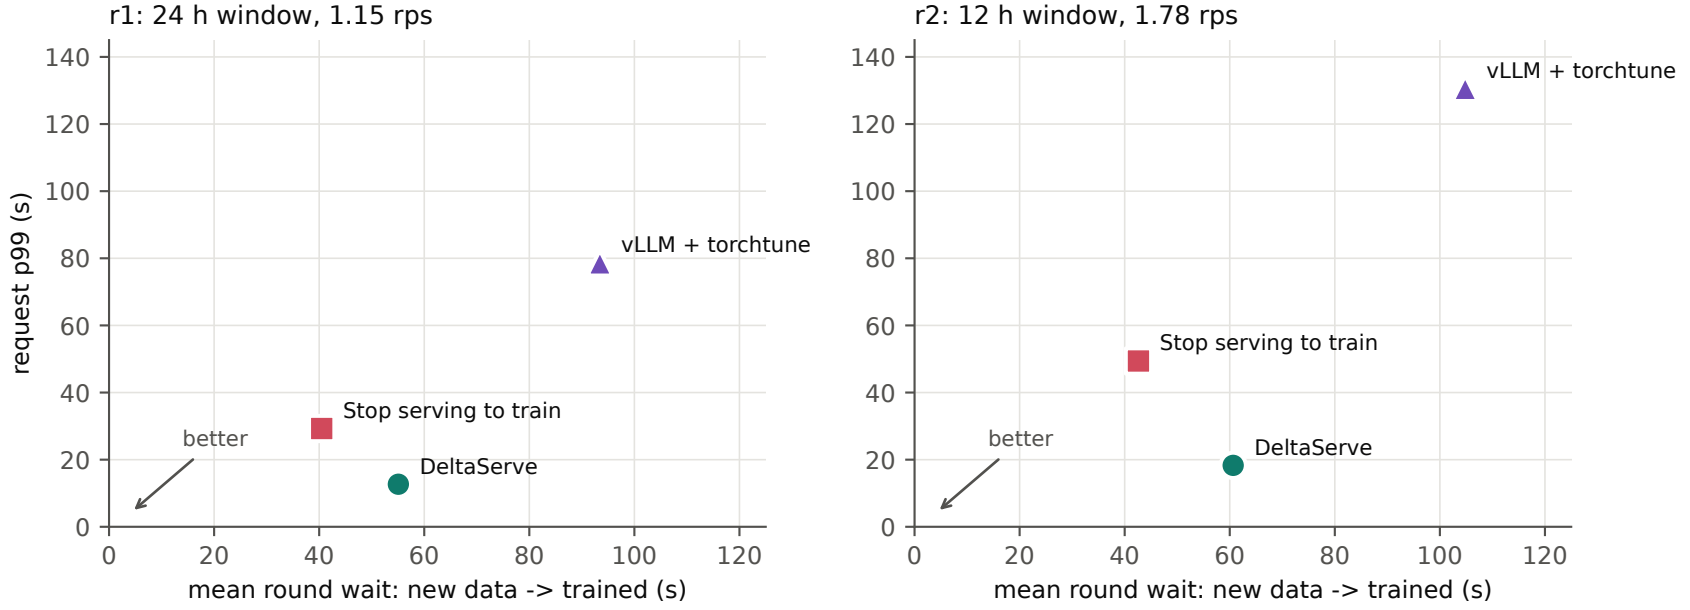
r1: 24 h window, 1.15 rps
request p99 (s)
mean round wait: new data -> trained (s)
vLLM + torchtune
Stop serving to train
DeltaServe
better
r2: 12 h window, 1.78 rps
mean round wait: new data -> trained (s)
vLLM + torchtune
Stop serving to train
DeltaServe
better